

## Fetch Dataset & Load Features

```
# Fetch Dataset...
# -----

# from ucimlrepo import fetch_ucirepo

# data_power_consumption = fetch_ucirepo(id=235)

# print(type(data_power_consumption))
# print(data_power_consumption)

✓ 16s

C:\Users\j...> cd C:\Data\Local\Programs\Python\Python311\Lib\site-packages\ucirepo\fetch.py:97: DtypeWarning: Columns (0, 1, 2, 3, 4, 5, 6, 7, 8) have mixed types. Specify dtype option on import or setting pandas option 'mode.infer_objects' to False.
df = pd.read_csv(data_url)

<class 'ucirepo.dotdict.dotdict'>
{'data': {'ids': None, 'features':
  Date      Time  Global_active_power  Global_reactive_power \
0  16/12/2006  17:24:00      4.216      0.418
1  16/12/2006  17:25:00      5.360      0.436
2  16/12/2006  17:26:00      5.374      0.498
3  16/12/2006  17:27:00      5.388      0.502
4  16/12/2006  17:28:00      3.666      0.528
...
2075254 26/11/2019  20:38:00      0.366      0.0
2075255 26/11/2019  20:39:00      0.944      0.0
2075256 26/11/2019  21:00:00      0.938      0.0
2075257 26/11/2019  21:01:00      0.934      0.0
2075258 26/11/2019  21:02:00      0.932      0.0

  Voltage  Global_intensity  Sub_metering_1  Sub_metering_2 \
0  234.840      18.400      0.000      1.000
1  233.630      23.000      0.000      1.000
2  233.290      23.000      0.000      2.000
3  233.740      23.000      0.000      1.000
4  235.680      15.800      0.000      1.000
...
2075254  240.43      4.0      0.0      0.0
2075255  240.8      4.0      0.0      0.0
2075256  239.82      3.8      0.0      0.0
2075257  239.7      3.8      0.0      0.0

  no
5   no
6   no
7   no
8   no
}

Output is truncated. View as a table or open in a text editor. Adjust cell output settings...
```

## Data after EDA

```
[1] ✓ 1m 51.3s

... C:\Users\j...> cd C:\Data\Local\Programs\Python\Python311\Lib\site-packages\ucirepo\fetch.py:97: DtypeWarning: Columns (0, 1, 2, 3, 4, 5, 6, 7, 8) have mixed types. Specify dtype option on import or setting pandas option 'mode.infer_objects' to False.
df = pd.read_csv(data_url)

Dataset Loaded Successfully!
Shape of X (features): (2049280, 9)
Shape of y (target): (2049280, 1)

Preview of Cleaned Data:
  Date      Time  Global_active_power  Global_reactive_power  Voltage \
0  16/12/2006  17:24:00      4.216      0.418      234.84
1  16/12/2006  17:25:00      5.360      0.436      233.63
2  16/12/2006  17:26:00      5.374      0.498      233.29
3  16/12/2006  17:27:00      5.388      0.502      233.74
4  16/12/2006  17:28:00      3.666      0.528      235.68

  Global_intensity  Sub_metering_1  Sub_metering_2  Sub_metering_3
0      18.400      0.000      1.000      17.0
1      23.000      0.000      1.000      16.0
2      23.000      0.000      2.000      17.0
3      23.000      0.000      1.000      17.0
4      15.800      0.000      1.000      17.0

  Global_active_power
0      4.216
1      5.360
2      5.374
3      5.388
4      3.666
```

## Model Evaluation Results:

```
Training Linear Regression...

Training Random Forest...
C:\Users\j...> cd C:\Data\Local\Programs\Python\Python311\Lib\site-packages\sklearn\base.py:139: DataConversionWarning: A column-vector has been converted to a Series; this may cause
return fit_method(estimator, *args, **kwargs)

Training Gradient Boosting...
C:\Users\j...> cd C:\Data\Local\Programs\Python\Python311\Lib\site-packages\sklearn\ensemble\_gb.py:87: DataConversionWarning:
y = column_or_1d(y, warn=True) # TODO: Is this still required?

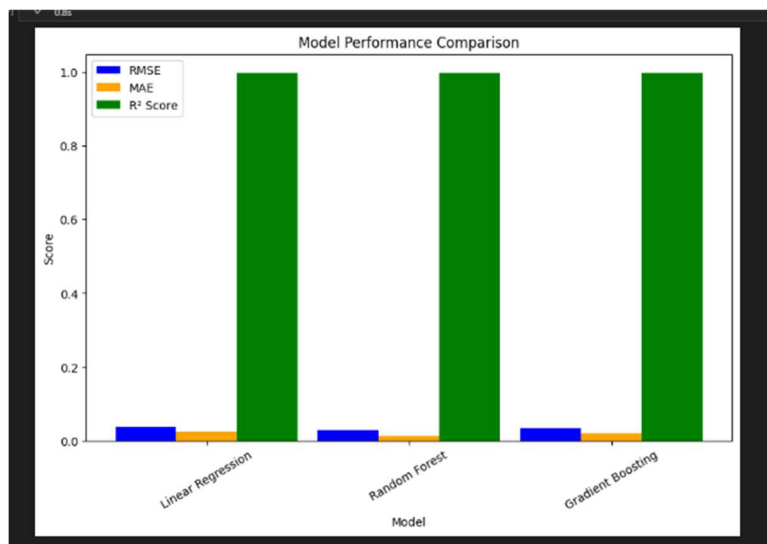
Model Evaluation Results:

--- Linear Regression ---
RMSE: 0.8384
MAE: 0.8248
R2: 0.9976

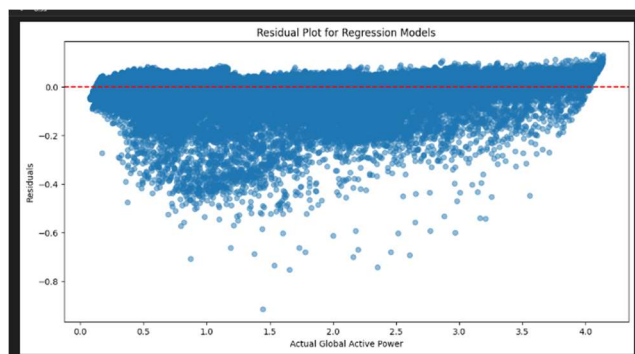
--- Random Forest ---
RMSE: 0.8287
MAE: 0.8126
R2: 0.9987

--- Gradient Boosting ---
RMSE: 0.8339
MAE: 0.8198
R2: 0.9981
```

Compare model performance and select the best-performing model.



Residual Plot



Model predictions

```
for model_name, model in models.items():
    y_pred = model.predict(X_test_scaled)
    print(f"{model_name} predictions: {y_pred[:5]}")  # Print first 5 predictions
```

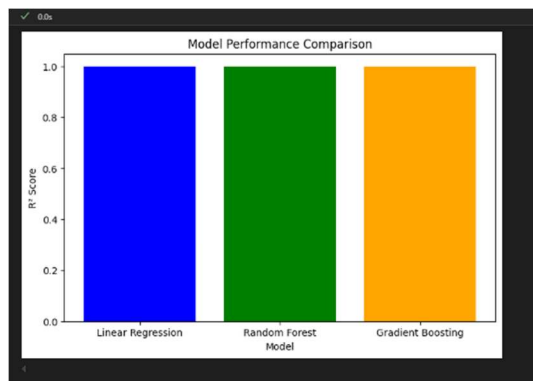
✓ 2m 8.3s

Linear Regression predictions: [[0.26157429]  
[1.6013208 ]  
[2.45168676]  
[0.39832884]  
[1.7481561 ]]

Random Forest predictions: [0.25268 1.64378 2.44156 0.3938 1.76324]

Gradient Boosting predictions: [0.26664603 1.61746032 2.45352508 0.39917085 1.7560915 ]

R<sup>2</sup> Score Bar Plot



## Classification Model

```
# Evaluate Performance
# -----
accuracy = accuracy_score(y_test, y_pred)
print("Classification Accuracy:", accuracy)
```

[26] ✓ 3m 38.2s

```
... Class distribution:
Global_active_power
Low      1923656
Medium   108077
High      17547
Name: count, dtype: int64
Classification Accuracy: 0.9970672626483448
```